

tensors and autograd

Week 10

- Lecture 1

- Why Deep learning
- The perceptron
 - activation function
- Multi-layer Perceptrons (a.k.a. Neural networks)

- Lecture 2


- tensors
- automatic differentiation (autograd)
- PyTorch for example

- Tutorial:

- An end-to-end example of using PyTorch to solve a non-linear regression problem.

Neural networks: linear maps plus simple non-linearities

- I. non-linearities provide the bare-minimum in expressive power. Then...
- II. more layers \square richer mappings possible
- III. more neurons per layer \square more flexibility




demo:

<https://playground.tensorflow.org>

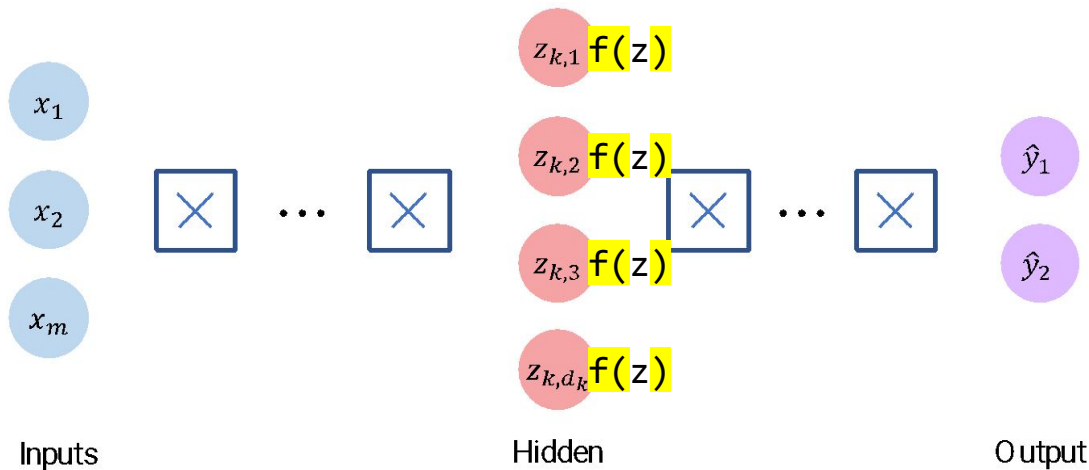
[g](#)


HOWEVER too much freedom can \square overfitting!

We attempt to
“regularise” with
penalties and other
constraints on model
power (“smoothing”)



Deep Neural Networks, Feed forward




a neural net is really an
alternating sequence of
matrix multiplies  and
**element-wise
non-linearities**

problems with gradients in really deep nets

Training a big neural net is computationally expensive (many epochs).

There are many “architectures” one might consider...

- number of layers
- number of nodes per layer




Depth was a major problem:

- gradients tend to (a) vanish and (b) get “shattered” in regular, deep, nets


Two particular techniques that seem to help:

BatchNorm (“batch normalisation”) - a scheme for “normalising” gradients all the way along.

Implemented in frameworks as just another kind of layer.



ResNets, with “skip” connections:



TensorFlow *versus* PyTorch

TensorFlow - developed at GoogleBrain (released 2015). Google used it for research and production. Intuitive high-level APIs such as Keras.

Static graph : user first defines the computation graph of model and then runs the ML model (compile
□ execute)

- more features, perhaps better for mobile and embedded deployments
- need "Sessions" and "Placeholders"...
- tensorboard – for visualising real-time accuracy graphs while training

PyTorch – developed at Facebook (released 2016). No API as it smoothly integrates with the python data science stack and is similar to NumPy.

Dynamic graph allows defining / manipulating the graph on the go. Imperative, interpreted, on-the-go.

- attracting Python developers
- no need to create session or placeholder objects
- scripting, just like using NumPy, easier to debug

Both essentially provide fast (GPU-enabled) **tensors** and **graphs**, including **autograd**. In 2023 it seems:

- **TensorFlow** with Keras : industry - make things faster and build ML products at scale
- **PyTorch**: research-oriented developers and python programmers - more customisable

We'd prefer you use PyTorch for the final project

PyTorch



Follow




I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

11:56 AM - 26 May 2017

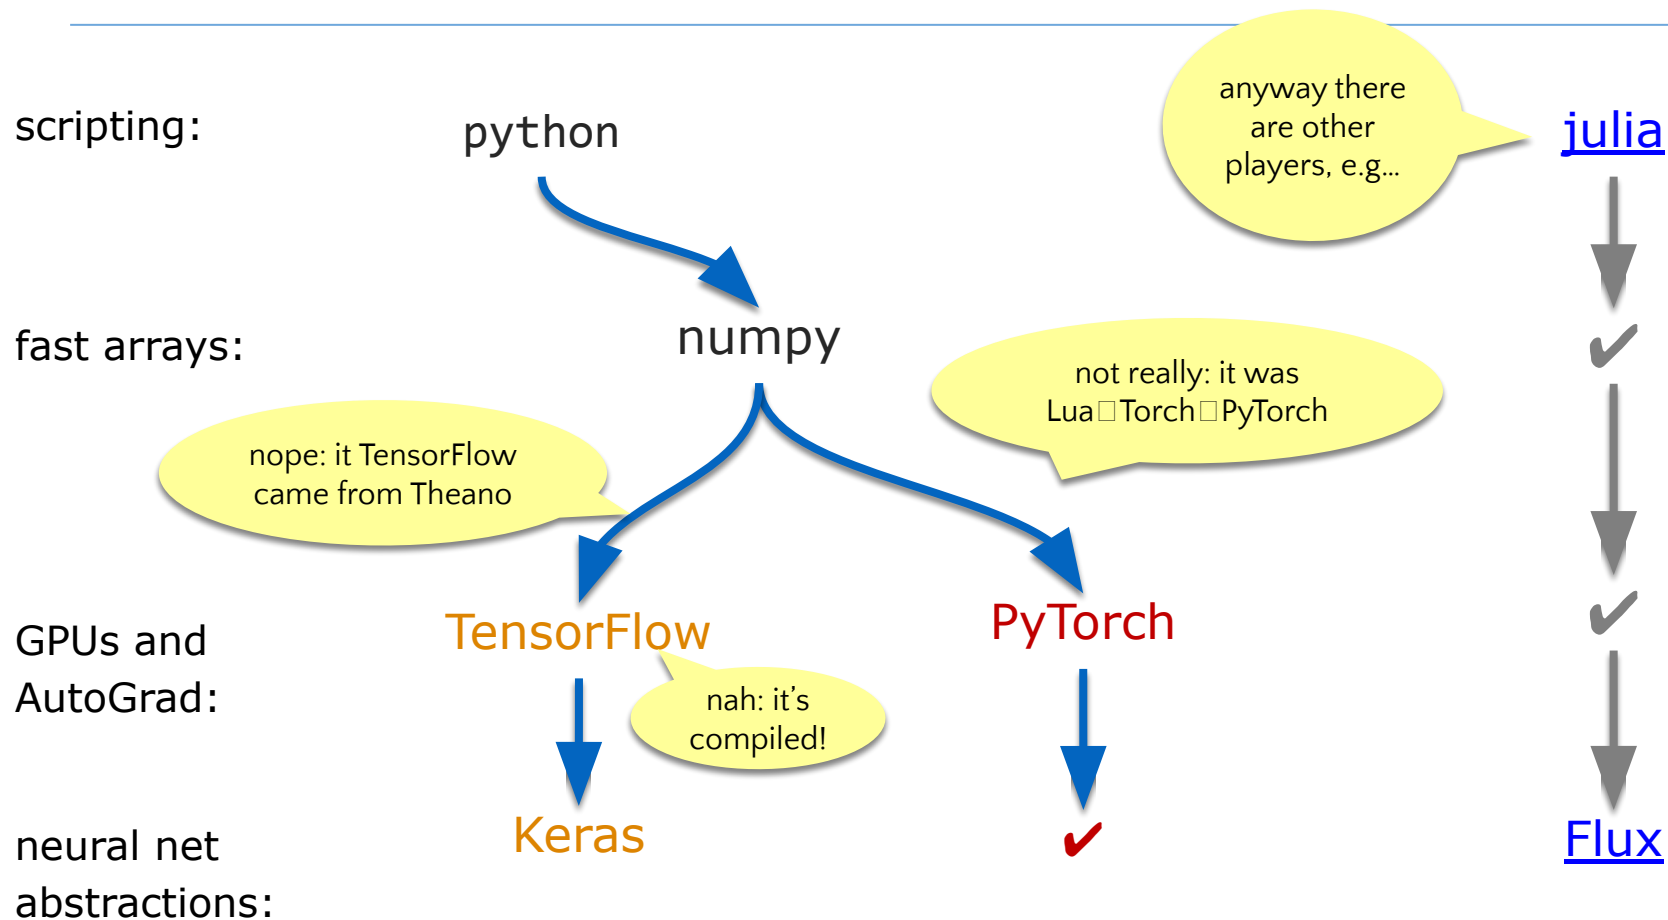
424 Retweets 1,706 Likes



33 424 1.7K




loose and bad history




PyTorch

Basically, PyTorch is like numpy but + **GPUs** and **autograd**

Essential data structure is the **tensor**




Scalar Vector Matrix Tensor



what we're *not* going to do, with functions

Function “compose”: the grandmother of *Mia* is *Beatrice*




And we can often go backwards via the `inverse` of the forwards function...

However `inverses` are *NOT*
what we're talking about
next!

gradients of a function

$$\frac{\partial x}{\partial x} = 1$$

$$L = v \times a$$



we can think of both the forward and the gradient computations as processing steps on a computational graph.

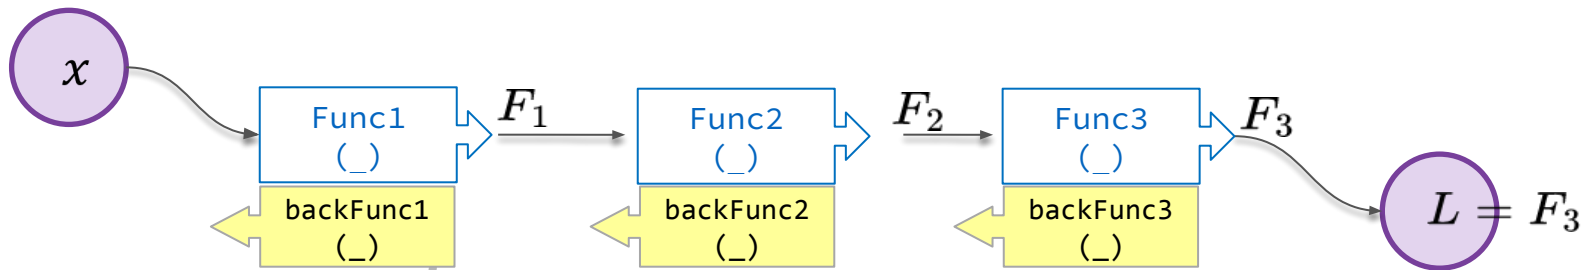
It makes sense to store those **gradient functions** at nodes on the same graph as the **forward** functions.

the chain rule

Say I have a function of a function of a function...

for example, $L = \text{Func3}(\text{Func2}(\text{Func1}(x)))$

To calculate L we do a forward pass through this graph. Each node remembers its forward-travelling values.



Q: how does L change if we wiggle x?

Ans: Chain Rule of calculus

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial F_3} \frac{\partial F_3}{\partial F_2} \frac{\partial F_2}{\partial F_1} \frac{\partial F_1}{\partial x}$$

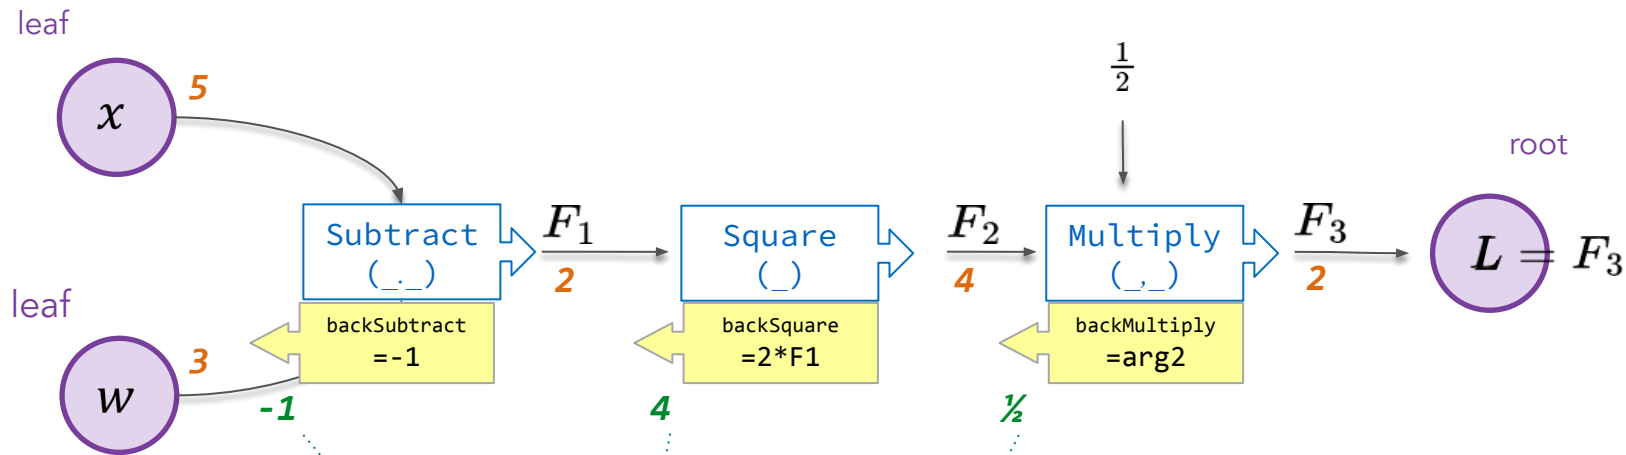
forwards we
compose,
backwards we
multiply



autograd: automatic differentiation

Consider finding the gradient of this function: $L = \frac{1}{2}(x - w)^2$

when the inputs were 5 and 3



Gradient for w at that point is $-1 * 4 * \frac{1}{2} = -2$

c.f. Chain Rule:

$$\begin{aligned}\frac{\partial L}{\partial w} &= 2 * \frac{1}{2} (x - w) * (-1) \\ &= -2\end{aligned}$$

“autograd”

- Implemented in TensorFlow and PyTorch (and elsewhere)

They “already know” `Back...()` for a set of mathematical operations `Forw()`, for example...

- ✓ plus, sum
- ✓ times (“Mul”)
- ✓ matrix multiplication (“MatMul”)
- ✓ exp, log,... and so on and on...
- ✓ batchNorm...
- ✓ linear...

Note for the younger user:
When I was a kid in grad school, we had to write our own, and it hurt, and we called it “backpropagation”

Autograd makes things so so *so much more pleasant*.

Regularisation – done better next week I think!....

discourage *overly complex* models, to improve generalization of our model on unseen data

- Regularization 1: Penalties on weights

- During training, pull a bit towards smaller values of the weights
- L1 vs L2 (minimize $\text{abs}(w)$ versus w^2)

- Regularization 2: Dropout

- During training, randomly set some activations to 0
- Typically 'drop' 50% of activations in layer
- Forces network to not rely on any one node

- Regularization 3: Early Stopping

- Stop training before we have a chance to overfit

Q: How to set regularisation parameters?...

