

AIML430/COMP309: ML Tools and Techniques

Lecture 10: Data Preprocessing

Ali Knott

School of Engineering and Computer Science, VUW



The plan for this week and next week

This week and next, we're finishing up with data preprocessing.

Today a bit more on a few things we've already touched on:

- Categorical data: how to *encode* it
- Numerical data
 - More on *scaling* (**normalisation** and **standardisation**)
 - More on *splitting* (more formally, **discretisation**).

Tomorrow: a more detailed look at *dimensionality reduction* methods.

- *Global methods* for changing feature space (PCA, & new ones)
- *Local methods* for removing features.

Next week (with Marcus Frean): 'Feature engineering'.

- In particular: how to *construct* useful new features.

Before we start - a few more points on data splitting!

Data splitting is when you divide your dataset into **training**, **validation** and **test** sets.

1. What proportions should you divide your dataset into?

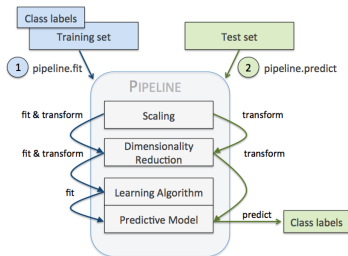
- No hard and fast rules. . .
- A reasonable range is
 - 60-80% training set
 - 10-20% validation set
 - 10-20% test set.

Before we start - a few more points on data splitting!

2. An important rule: *split first*, and then preprocess the separate datasets separately!

- If you preprocess *before* splitting, information from the training set can *leak* into the validation/test sets.
- That **data leakage** compromises the validation/test sets.

This hygiene is built into Scikit-learn's pipeline structures. . .




Before we start - a few more points on data splitting!

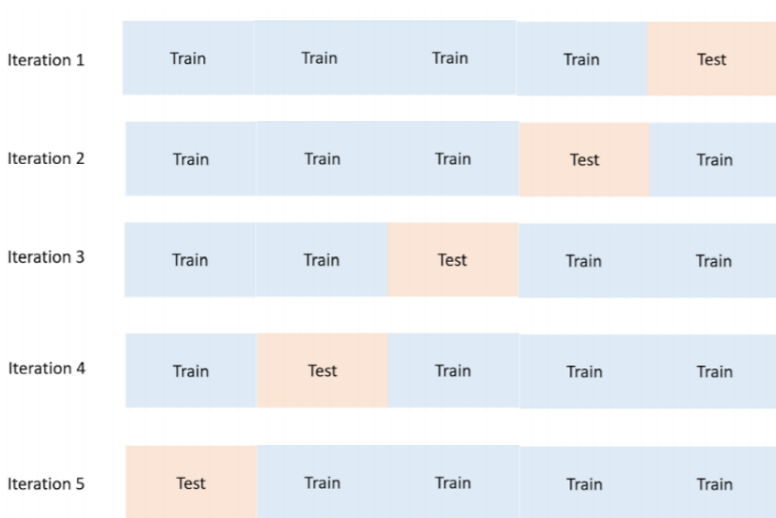
3. We briefly mentioned a useful method called **cross-validation**: here's some more on that.

If you don't have much data, leaving some out for testing feels wasteful.

- **k -fold cross-validation** is a way of training on all the data.
- Method: split the dataset into k equal-sized **folds**.
 - For each fold F :
 - Set aside F for testing, and train on all the other folds.
- To evaluate the learned model, take the *average* test performance.



An example: 5-fold cross-validation



Cross-validation notes

There's no 'leakage' in cross-validation, because each model trains separately.

But: there's no real *validation set* in the method just sketched.


- If you want to tune model parameters, you can hold one fold aside as a true 'test set', and do cross-validation (with tuning) on the remaining $k - 1$ folds.

One more point on training/testing... about overfitting

4. Some training processes take time, and increase the complexity of the learned model over time.
- A good example is a neural network.
 - This trains incrementally, over n iterations...
 - Weights are gradually adapted to training items as time passes.

In this kind of case, *overfitting happens at a point in time*.

- To identify that point, test on the validation set regularly *during training*, and see when validation performance *stops improving*.



And now to data preprocessing!

Recap: there are two broad types of feature:

- **Categorical** features (with discrete alternative values)
- **Numerical** features (with continuous values).

For categorical data, I'll say more about how it should be *encoded*.


For numerical data, I'll say more about how it should be *scaled*, and *discretised*.

1. Encoding categorical data

1.1. 'One-hot' encoding

When we introduced neural network classifiers, we focussed on **binary classifiers**, that make decisions about a *single class*.


- 1 means 'yes', 0 means 'no'.
- We saw a classifier that responds discretely, and one that responds with a *probability*.



1.1. 'One-hot' encoding

Neural networks can also be *multi-class* classifiers.

- If there are n classes, the network has n output units.
- Here's an image classifier, that recognises discrete object classes.
- Activity in the output layer can be interpreted as a probability distribution—if we normalise it to sum to 1.



1.1. 'One-hot' encoding

This kind of network is **trained** on a **one-hot** probability distribution: 1 for the correct class, and 0 for all the others.

- In the dataset, the output feature must be *encoded* in this one-hot notation.
- We often do that by creating a **dummy variable** for each output class.


Country	dummy variables		
	USA	UK	France
USA	1	0	0
UK	0	1	0
USA	1	0	0
France	0	0	1
USA	1	0	0
UK	0	1	0

1.2. Ordinal encoding

Some categorical features are **unordered**; others are **ordinal**.

- An **ordinal** feature has values with a natural *ranking*:

- E.g. *poor, good, very good, excellent*



An **ordinal encoding** just assigns *integers* to the ranked categories.

Original Encoding	Ordinal Encoding
Poor	1
Good	2
Very Good	3
Excellent	4

2. Encoding numerical data

Two topics here

- Scaling
- Discretisation.

2.1. Scaling numerical data

We have already discussed how numerical features should be *scaled*:

- If you don't scale, features with a larger *range* will dominate. . .
- Especially in *distance* computations.

I'll discuss two scaling methods in some more detail.

- Linear scaling
- Standardisation.

2.1.1. Linear scaling

In **linear scaling** for a feature f , we scale the whole dataset by a constant factor.

Often, we scale so that datapoints fall in the range $[0,1]$.

- We begin by finding the *maximum* and *minimum* values for f (x_{max} , x_{min}).
- Then, for any value x , the scaled value x' is given by $\frac{x - x_{min}}{x_{max} - x_{min}}$.
- In this scheme, the min values are scaled to 0, and the max values are scaled to 1.

We can also scale to a specified range, $[New_{min}, New_{max}]$. In this case,

$$x' = \frac{x - x_{min}}{x_{max} - x_{min}} \times (New_{max} - New_{min}) + New_{min}.$$

In scikit-learn, the **MinMaxScaler** object does these operations.

2.1.2. z-score standardisation

If a feature happens to be normally distributed, there's a more informative way of scaling.

z-score standardisation, or **centre scaling**, scales the *distribution* of points, so they have a *mean of 0*, and a *standard deviation of 1*.

- This kind of scaling is useful, because its units are standard deviations. (Value 1 is 1 standard deviation above the mean.)
- Values map easily to probabilities (see next page).

For a feature with mean μ and standard deviation σ , the scaled value z for a raw value x is given by


$$z = \frac{x - \mu}{\sigma}$$

To apply this kind of scaling, use the scikit-learn object **StandardScaler**.

z-scores and p values

With a normally distributed variable, we can compute the probability that values fall within a given range.

- The p value for a given range is the area under the normal curve for that range. There's a mapping from z ranges to p values.
- You can compute p values with the `scipy.stats` fn `norm.ppf`.



Linear scaling or z-score standardisation?

The basic rule:

- If your numerical feature is normally distributed, use z-score standardisation.
- Otherwise, use linear scaling.

If you have a *mixture* of normal and non-normal input features, you should use linear scaling on all of them—so you can scale them all to *the same range*.

2.2. Discretisation

Discretisation (or **binning**) is the process of converting continuous numeric values (e.g. price, age, weight) into discrete intervals.

Two key purposes for discretisation:

- It's required for some ML algorithms—in particular, decision trees.
- It's used in many visualisation routines.

There are two types of discretisation:

- **Supervised**—splitting is guided by class labels. (E.g. in decision tree learning.)
- **Unsupervised**—no reference to class labels.


2.2.1. Unsupervised discretisation methods

Uniform (equal-width) discretisation creates equal-sized **bins** for a numerical variable.

- This is what we use to make a **histogram**.
- The basic process: find the Max and Min values, then divide Max-Min by the desired number of bins.

Say we have this set of temperature values, and we want 7 bins.

- 85, 80, 83, 70, 64, 65, 68, 71, 69, 72, 75, 75, 81, 72
- The width of each bin $(85-64)/7=3$.
- `KBinsDiscretizer(n_bin=7, encode='ordinal', strategy='uniform')`




2.2.1. Unsupervised discretisation methods

Quantile (equal-depth) discretisation divides numerical data into bins that each hold (approximately) the same number of items.

- `KBinsDiscretizer(n_bin=4, encode='ordinal', strategy='quantile')`

For the temperature values shown before, we get:



I don't see what this is useful for!


Useful to note: discretisation creates an **ordinal** categorical variable.

2.2.1. Unsupervised discretisation methods

Clustering can be used to choose bins for discretisation.

- `KBinsDiscretizer(n_bin=3, encode='ordinal', strategy='kmeans')`
- This runs *k*-means clustering *in one dimension*, with three centroids.
 - A bin is created for each cluster.

For the temperature values shown before, we get:



2.2.2. Supervised discretisation methods

Supervised discretisation methods aim to find splits for a numerical variable that best match values of a discrete output class.

The basic idea is to search through *all possible split points*, to find the points that best separate classes.

- Often, we look for *binary* splits, that distinguish one class from the others.
- Within each binary split we must *recurse*, looking for further splits:
 - To separate other classes;
 - To refine splits for each class.

The process is similar to decision-tree learning.

2.2.2. Supervised discretisation methods

There are various methods for assessing the ‘goodness’ of a candidate split.

Some methods compute the **total error** of a given split.


- For each split, assign the most frequent class label. . .
- Then count the number of items that are of other classes
- And sum over all splits.
- This is used by **1R discretisation**.
 - In **R**’s **dprep** package, use the function **disc.1r**.

2.2.2. Supervised discretisation methods

Other supervised methods compute the **entropy** (or ‘**purity**’) of a given split.

Entropy is a probabilistic measure: it quantifies the ‘sharpness’ of a probability distribution. (A sharp distribution has low entropy.)

- For a probability distribution over N classes $c_1 \dots c_N$, the entropy is given by $E = - \sum_{i=1}^N p_c \times \log_2(p_c)$.
- For each candidate split, we can estimate probabilities by counting, and compute entropies. . .



- We are looking for the split with the *lowest entropy*.
- In R's **dprep** package, use the function **disc.mentr**.

Temporary page!

\LaTeX was unable to guess the total number of pages correctly, so
there was some unprocessed data that should have been added to the
final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will be
removed, because \LaTeX now knows how many pages to expect for the
document.