

# AIML430/COMP309: ML Tools and Techniques

## Lecture 11: Dimensionality Reduction Methods

Ali Knott

School of Engineering and Computer Science, VUW



# The plan for today and next week

Today (and next week) we're looking at 'feature manipulation' methods.

Today: a more detailed look at *dimensionality reduction*.

- *Global methods* for changing feature space (PCA, & new ones)
  - Sometimes called **feature extraction**.
- *Local methods* for removing features
  - Sometimes called **feature selection**.

Next week (with Marcus Frean): 'Feature engineering'.

- In particular: how to *construct* useful new features.
  - Discretisation (yesterday) is one way to do this. . .
  - Marcus will show you others!

# Why do feature manipulation?

Raw data contains all sorts of features. . .

- Not all necessarily useful for your purposes.

Some reasons for manipulating features:

- To improve performance of the model to be learned.
  - Improvements could be in speed, predictive performance, simplicity, interpretability. . .
- To reduce noise. (And to reduce overfitting.)
- To compress the data.
- To visualize something in the data.

# 1. Feature extraction

Feature extraction is a *global* transformation of the feature space.

- We start with a space of  $d$  dimensions, and transform to  $m$  dimensions. ( $m$  normally smaller than  $d$ .)
- The idea: to retain most of the relevant information in the original feature space.

# Feature extraction

There are two main approaches to feature extraction.

One is **projection**: project points into a lower-dimensional space.

- To visualise projection: think of points ‘falling onto’ a straight line, or plane.

The other is **manifold learning**.


- A manifold is a **topological space** that maps to your feature space.
- The manifold can be *curved*—so we can model complex patterns.
- But at any *local* point, its topology *resembles* that of the feature space.
  - Your data points might live on a *sphere*, or a *saddle*, defined within your feature space. . .

# 1.1 Projection techniques for feature extraction

We have already seen two projection techniques: **principal component analysis (PCA)** and **independent components analysis (ICA)**.


Another one is **linear discriminant analysis (LDA)**.

- This method simplifies *input space* for a classification task.
- It finds linear decision boundaries in feature space, and projects data onto planes perpendicular to those boundaries.
- `sklearn.discriminant_analysis.LinearDiscriminantAnalysis`



## 1.2 Manifold techniques for feature extraction

A useful visualisation of reducing dimensionality with a manifold:




We already mentioned **t-sne**.

- The idea here is to transform feature space into an **embedding space**. . .
  - With fewer dimensions. . .
  - Which is *optimised*, to preserve the pairwise distances between all datapoints.
- Optimisation happens iteratively, using gradient descent.

## 1.2 Manifold techniques for feature extraction

We've also already seen **autoencoders**.

- This is a multi-layer perceptron that learns to map each datapoint *back onto itself* . .
- Through a 'bottleneck layer' *smaller* than the input/output layers.
- '**Embedding**' is a ML term used for any learned representation.





## 2. Feature selection techniques

The idea here is to *remove individual features* from the original dataset.

- Find uninformative or redundant features, and get rid of them!

There are two types of FS technique.

- **Univariate** methods consider each feature by itself.
  - Mostly: how much does it *covary* with other features?
- **Multivariate** methods consider all features at once.
  - E.g. in algorithms that iterate through each feature.

## 2.1. Univariate feature selection methods

Univariate methods run *tests* on each input feature, to see how it relates to other features.

- Strong relationships with the *output feature* are good!
- Strong relationships with *other input features* are bad!

Based on the results of these tests, you can em select features.

- Keep the good ones, discard the bad ones!


We'll review some of the tests you can use.

- All in sklearn's **feature\_selection** module.
- Key functions: **SelectKBest**, **SelectPercentile**


## 2.1. Univariate feature selection methods

For a *regression* task, you should use tests based on *linear correlation*.


**Positive correlation**



**Negative correlation**



**No correlation**



Use `sklearn.feature_selection.r_regression()`

## 2.1. Univariate feature selection methods

For a *classification* task, you should use tests based on a *discrete* output variable.

- A good test is **chi-squared**: this test identifies features that are likely to be independent of the output class.
- It tests the (null) hypothesis that there's no association between a feature and the output feature.

Use `sklearn.feature_selection.chi2()`

# Sklearn offers several other univariate tests...

There are methods based on **mutual information** calculations.

- Mutual information measures the *reduction in uncertainty* about Variable 1, that comes from knowing the value of Variable 2.
  - For regression tasks, use `sklearn.feature_selection.mutual_info_regression`
  - For classification, use `sklearn.feature_selection.mutual_info_classif`

For continuous variables that aren't normally distributed, Spearman's rank **correlation coefficient** is a good measure of correlation.

- Use `scipy.stats.spearmanr`

For continuous inputs and discrete outputs, **Analysis of Variance (ANOVA)** is also useful. (Use `sklearn.feature_selection.f_classif`)

- This finds the amount of variance in the output class that's *explained* by the input variable. (Expressed as a proportion.)

## 2.1. Univariate model-based feature selection

Another way of identifying (individual) features is to train a model, and then *analyse the model*.

- This is done with `sklearn.feature_selection.SelectFromModel`.
- With a *regression* model, for instance, scores can come straight from *coefficients* for input variables.
- With a *decision tree*, there are higher scores for the features used high up in the tree.

## 2.2. Multivariate feature selection methods

Selecting individual features is not always ideal, because sometimes features *interact*: so their importance can't be evaluated in isolation.

- **Multivariate feature selection** methods assess the importance of *groups of features*.

Multivariate methods are often implemented as 'wrappers' around model training processes.

- The basic idea is to run *many* training processes, experimenting with different features—and see which feature set does best.
- Obviously, it's *costly* to iterate many times over the whole training process!

## 2.2. Multivariate feature selection methods

There are two ‘standard’ wrapper methods for feature selection.

**Sequential Forward Generation (SFG)** starts with an empty set of ‘used features’: all features are ‘candidates’. At each iteration:

- For each candidate feature  $C$ , train a model that adds  $C$  to the set of used features. Pick the best-performing model.

This works best when the optimal feature set is small. (But it’s **greedy**!)

**Sequential Backward Generation (SBG)** starts with a full set of ‘used features’. At each iteration:

- For each used feature  $U$ , train a model *without*  $U$ . Pick the model whose performance decreases least.

This works best when the optimal feature set is big. (It’s also greedy!)



## 2.2. Multivariate feature selection methods

Advanced feature selection methods take less 'greedy' approaches to selecting/rejecting features.

- Essentially, we're searching the space of *possible feature combinations*.
- If there are  $n$  features, there are  $2^n$  possible combinations...  
How can we effectively explore that space?

For large feature sets, *exhaustive* search is *impossible*.

- We have to find good *heuristic* search methods.
- Our school specialises in this work!

A common heuristic search method is with **genetic algorithms**.

# Genetic algorithms: A tiny taster

Imagine a population of ‘creatures’ that have *feature sets instead of genes*.

- They compete with each other, and only the fittest survive. . .
- Fitness is determined by how well their model performs!
- The survivors are randomly paired together, to ‘breed’ a *new generation* of creatures.

If you fancy this kind of synthetic biology, Vic’s AI group is for you! :-)

# Summary: 'Dimensionality reduction' methods

*Global methods* for dimensionality reduction map data into *new feature spaces*.

- **Projection** methods (PCA, ICA, LDA) find good *linear* mappings. . .
- **Manifold** methods (t-SNE, autoencoders) learn good *nonlinear* mappings.

*Local methods* for dimensionality reduction *select* features from the full feature set.

- **Univariate** methods score *individual* features, & keep the best. . .
- **Multivariate** ('wrapper-based') methods assess *feature combinations*.
  - Simple methods (forward and backward) perform 'greedy' searches of feature combinations.
  - 'Heuristic' searches (e.g. genetic algorithms) are smarter—but also very expensive.

# ... And now over to Marcus Frean!

Next week, Marcus will take over.

The first thing he'll look at will be 'feature engineering': that is, how to *construct new features*.

- Discretisation is a simple way to do that. . . he will show you others!